

Yonah Citron

London · Dual UK/US Work Authorisation
citronyonah@gmail.com · yonahcitron.com · github.com/yonahcitron

Summary

Cambridge linguistics graduate with 3.5 years of production software experience, now completing an MSc at Imperial College London focused on systems programming, compiler infrastructure (LLVM), and formal verification. Scoped and delivered data-intensive platforms, ETL and NLP pipelines, API backends, and web applications in Python, C#, JavaScript, and SQL. Experienced in owning projects end-to-end and deploying to cloud using IaC and CI/CD.

Skills

Languages: Python, C++, C#, Bash, SQL, JavaScript. **Systems:** LLVM toolchain, KLEE symbolic execution, OS internals, concurrency. **Data:** PostgreSQL, Elasticsearch, MongoDB, Databricks. **Cloud & DevOps:** Azure, AWS; Bicep, Terraform, Kubernetes; GitHub Actions, Azure Pipelines. **AI/NLP:** LLM orchestration, RAG, vector search, Langchain/Langgraph.

Education

MSc Computing · Imperial College London *expected Oct 2026*

2025 – 2026

- Electives: Advanced Computer Architecture; Networks & Distributed Systems; Computer Vision; Graphics; Mathematics & Logic
- Core (C++): Operating Systems & Concurrency; Algorithms & Data Structures; Software Engineering
- Research: Designed and built pSWIG, a parser test-input generator combining incremental symbolic execution with coverage-guided and keyword-directed search

MA, BA Linguistics · University of Cambridge *First Class Honours*

2017 – 2021

- David Thompson Prize for outstanding academic performance in final year
- Dissertation applying data-analytic techniques to a Neo-Aramaic corpus grammar to explore the effects of syntactic frequency and collocation

Work

Shell *initially via Kubrick Group, then permanent*

Jan 2022 – Oct 2025

Threat Intelligence Platform & Console

- Built full-stack platform ingesting, normalising, and enriching external threat feeds into a centralised data warehouse, with a web console for real-time search, classification, and triage
- Designed strongly-typed data interface for consolidating diverse feed schemas; implemented on-demand NLP pipelines for translation and summarisation
- Implemented transparent auditing logs and data-state tracking
- Deployed infrastructure across staged environments using parameterised IaC and CI/CD

Distributed Document Processing Platform

- Core developer on a distributed .NET Core platform processing millions of internal documents; architected batch-processing improvements
- Automated legacy manual deployments with CI/CD pipelines, improving reliability and deployment velocity
- Documented microservice architecture to improve audit trail and onboarding speed

EV Charging Data & AI Systems

- Built LLM pipeline extracting and classifying work order data from unstructured emails into a Databricks warehouse with versioned schema migrations (Alembic)
- Implemented Medallion architecture for document ingestion and serving; built vector search endpoints and RAG models registered via MLflow
- Deployed continuous delivery pipeline with automated LLM evaluation to maintain answer quality over time

Competitive Intelligence Portal

- Built data aggregation and search portal consolidating multiple external intelligence feeds, with NLP classification, semantic search, and summarisation
- Optimised search and result-ranking in PostgreSQL using stored procedures; implemented semantic search over vector databases
- Built back-end APIs in Express.js, orchestrating data batches with composable, Dockerised pipeline stages and CI/CD
- Nominated for internal Shell award for exceptional contributions to project delivery

LLM Tooling & Legacy Modernisation

- Developed internal LLM framework for structured extraction from unstructured text across multiple teams; experimented with open-source models via llama.cpp on private infrastructure
- Designed custom pre-processing pipelines, including a graph-based email deduplication algorithm to reduce token usage
- Modernised legacy knowledge management app with RAG-based semantic queries and infrastructure-as-code deployments

Interests

Languages: Advanced Russian; intermediate Arabic and Hebrew. Fond of Tolstoy and Russian-language literature. **Music:** Choral singer (two choirs). **Hardware:** Build electric guitars and custom mechanical keyboards. **Computing:** Perpetually refining my Doom Emacs config, Arch Linux desktop, and other dotfiles.